

Discoveries

Original findings from a grown spiking creature, with no external antecedent

Arne Diegenbach

The aibaby project — a C++17 spiking neural creature grown from a DNA/genome file

8 October 2026 (revised)

Abstract

This document collects the findings of the aibaby project that are not held up by any published result. The project grows a spiking neural “creature” from a genome file and teaches it to produce vowel formants; much of its machinery is named after the literature it came from, and those mechanisms are deliberately excluded here. What remains is the set of things the system itself was the first to say: measured quantities with no external reference, structural obstacles that only appear in this architecture, and a group of inference failures that were found by being committed and then caught. The longest single thread is §4: because the readout is a rate-weighted centroid, the influence of one neuron on behaviour can be differentiated rather than estimated, and that derivative — taken with a sign, which is the part the data forced — accounts for a measured asymmetry, its reversal between targets, and why one restriction on plasticity costs half the lesson while another binds a third of all synapses and costs nothing.

The methodological findings are listed first and at length, because they are the most transferable part and the least specific to a creature made of spikes. Several of them are reports of a mistake that survived review — a test no outcome could pass, a refusal criterion computed from the wrong arm, a statistic diagnosed in a comment and then quoted for months. They are included on the view that a failure mode with a worked example is more useful than a rule without one.

Every quantity below comes from a run gated on a pinned determinism hash, `ad96f882becbee92`, under which every mechanism that ships off is bit-identical. Section 8 gives provenance.

1. Scope, and what is excluded

A finding is included if the project discovered it, and excluded if a published result says it first or frames it. The line is drawn at the *claim*, not at the tools: the estimators used throughout are Cochran’s ratio estimator and the Quenouille–Tukey jackknife, and the quantities they produce are still this project’s, so they appear here while the methods do not.

Excluded on that basis are: the account of interference as failed contextual inference (Heald, Lengyel and Wolpert’s COIN model), the finding that context gating alone is insufficient without synaptic stabilisation (Masse, Grant and Freedman), importance-weighted protection of earlier weights (Kirkpatrick et al.; Zenke, Poole and Ganguli), node perturbation as a motor learning rule (Fiete, Fee and Seung), the stability–plasticity dilemma (Grossberg), the uniform cluster prior implicit in MacQueen’s 1/wins rule, the division of a population code into output-potent and output-null dimensions (Kaufman, Churchland, Ryu and

Shenoy), and activity-dependent myelination as a consolidation mechanism (the oligodendroglia literature, cited by title and DOI in the main list where author lists could not be verified). Each is cited in the project's main reference list. Appendix A lists the project results that were excluded by this criterion, so the exclusion can be audited rather than taken on trust.

2. Measurement and inference

These came out of running the same experiment many times and being wrong in a way that printed a number. None has an external source; each has a worked example in the project history.

2.1 Chance has no fixed value when the readout is “nearest”

A k-alternative naming score read against $1/k$ is not a chance comparison. When the decision rule is nearest-target in formant space, chance depends on the *geometry* of the targets: the same uninformative voice scores 0.161 against targets on the vowel corners and 0.306 against a 2×2 arrangement. An informationless voice reads about 0.30 on a four-word task where $1/k$ would predict 0.25. One earlier ceiling claim was a null error of exactly this kind: a real effect scored against an arithmetic chance level while the matched control sat in the same table, unread.

2.2 A test that nearly all outcomes pass is not a test, and neither is one that none can pass

Two symmetric failures, both found after they had produced a published-looking verdict. A guard written as $cd2 > -3 \cdot 13.1$ — “greater than -39.3 ” — passes on nearly any data, and was reported as “F2 is spared” when it meant “F2 is unmeasured”. In the other direction, an intervention's success branch required a shift of 0.090 where a complete closure of the gap could produce only 0.070; it printed a confident conclusion from an unreachable branch. The fix in both cases is to print the power of the test beside its result, so that a criterion nothing could have cleared is visible rather than silent.

2.3 The reference point is the hypothesis's own prediction, not zero

A hypothesis that names a value must be tested against that value. A drift account predicting a decay of -0.034 was first tested for being distinguishable from *zero*, which it was not, and so survived; tested against its own prediction, the measurement sat 3.9 SE away and the account was excluded. The same error in another guise is scoring an effect against $1/k$ while the real control is in the table.

2.4 A null and a reversal are different results

An else branch that catches both reports the wrong one. An intervention expected to close a gap was tested only for closure; when the gap *widened*, the outcome fell through to the null branch and was written up as “no effect”. The widening was 2.5 SE in the opposite direction to the prediction, which is a stronger result than the null it was reported as.

2.5 A constant copied instead of derived fails silently, and travels

Six instances in one 34,000-line file, in two families. *Indices*: a literal array size that was a copy of a named bound and overran into the next field; positional grid indices; positional arm indices that would have

mislabelled whole rows; and pair indices ($k \cdot 2$, $k \cdot 2 + 1$) that silently shifted when two table rows were inserted above them. *Labels*: a band value copied beside the table it came from, so an experiment reported its widest condition as a tenth of its real width after the table was rescaled; and a deleted control whose label and refusal criterion were inherited by unrelated rows; and a single hardcoded baseline arm, correct while every arm shared one lesson target and silently wrong the moment a second target was added, which charged the new arms the whole difference between the two conditions — one cost read 11% against the old control and 3% against its own, and three of eight printed numbers were wrong.

What makes this class survive review is specific and worth stating: **every site carried a comment acknowledging the indexing was positional, and the comment was treated as the fix.** One even warned about a *skipped* row reading zero — a different hazard from a *shifted* row, against which the guard written for the first does nothing. A mislabelled number still prints with a standard error and still reads as a result, so nothing in the output looks wrong. The question that finds these is: *is this label attached to the data, or merely next to it?*

2.6 A passing determinism hash does not say the code is correct

The project gates every run on a pinned hash. A pass says the shipped configuration is unchanged and the run is reproducible; it says nothing about whether the experiment computes what its name claims. Four of the six errors in §2.5 were introduced and committed under a passing hash.

2.7 Reproduce the conditions, not the experiment name

A headline was called irreproducible after three runs at the registered default tick count. The recorded result had been taken at $3.5\times$ that budget, which was written in its own log's filename and in the note that cited it. A registered default is not evidence about what a recorded run used, and this system has a documented power curve across exactly that range — one experiment reads +1.0 at 560k ticks against +18.3 at 3.4M.

The same error in its archive form, and it is the easier one to commit. A ladder of mask extents was assembled by collecting arms from every log in the results directory with a single `grep -h` matching on the arm's *name*. One of the seven rows came from a pilot run whose own baseline arm gained +0.0363 where every other row's baseline gained +0.3618 — a tenfold difference in the quantity being compared — and nothing in the pooled output recorded which log each line came from. **A search across runs is a join on the arm name, and the arm name does not carry the budget.** The rule that catches it is to read the anchor of the log a number comes from before reading the number, which this project applies to its own runs as a matter of course and had not applied to its own archive.

2.8 Diagnosing a broken statistic does not retire it

A retention measure was computed as a mean of per-creature ratios whose denominator — how much that creature learned — was floored only at $1e-6$. The code carried a comment beside it recording that the statistic had read -5.587 ± 6.729 and 27.222 ± 21.160 on another genome and calling it “a broken statistic reporting nothing”, and a denominator-free replacement was added alongside. The headline, the milestone gate and the pre-registered specificity criterion all went on reading the broken one.

Measured on the same 128 sessions, the two estimators disagree in *sign*: a mean of ratios gives

$+0.182 \pm 0.663$ and a jackknifed ratio of sums gives -0.211 ± 0.074 , with standard errors differing ninefold. The transferable part is not the statistic but the order of operations: a comment records a fact, it does not change what the gates compute, and every consumer has to be moved by hand. Grep for the variable, not for the comment.

2.9 Unanimous signs on three seeds are not power

Repeatedly, effects that were unanimous and large on a handful of creatures did not survive more of them. A smoothing sweep read $+0.24$ at $n = 3$ and $+0.03$ at $n = 6$; a salience-reward effect was $+2.6$ SE at $n = 12$ and null at $n = 16$; a vigilance score of 0.737 fell to 0.594 at 96 creatures. The project adopted a 3 SE bar at $n \leq 18$ as a direct consequence, and the bar must be fixed once and named, not chosen per test: one verdict printed “not yet a finding” for 2.9 SE against a 3 SE bar and then made a 2.7 SE contrast its headline because that test had been coded against 2.0.

2.10 Two points agreeing to a tenth of a percent is one coincidence

Two lesson targets gave lever costs of 30.01% and 22.12% of the lesson, and the distance each had to move the readout stood in a ratio of 1.356 against the costs’ 1.357 — a match to 0.1%, from no fitted parameter. It was recorded as a coincidence on two points rather than a law, and the third target refused it: the law predicted 14.8% and the measurement was $25.0\% \pm 5.9\%$, and the three points are *non-monotone* in distance travelled ($0.200 \rightarrow 25.0\%$, $0.300 \rightarrow 22.1\%$, $0.407 \rightarrow 30.0\%$), which no proportionality can produce at all.

The transferable part is the arithmetic rather than the caution. Any smooth one-parameter family can be made to pass through two points, so a two-point agreement carries no information about the family; and the cheapest way to settle it was to measure a third point, which cost one overnight run, rather than to argue about the first two. The same project had already paid for the opposite choice once: a leverage law fitted to three costs with one free parameter matched them to 1.68 percentage points and then failed two further costs out of sample. Restricted to the points where the lever is the same neuron, the surviving pattern needs a large distance-independent term plus a weak distance term — two parameters on two points — so the account was stopped there instead.

2.11 A pre-registered prediction can be wrong in sign on both of its axes

Pre-registration fixes what a result will be compared against; it does not make the prediction right, and the failure mode worth recording is the one where a prediction fails in the direction nobody wrote a branch for. A cross-mechanism regularity — every plasticity restriction tried here cut learning more than it cut forgetting — was used to predict that removing a restriction would buy more learning than it cost in retention. The arm that removed it gained *less* and retained *less*. Both differences are nulls, so the regularity is not refuted at significance, but neither was it confirmed, and the registered text had no reading for “worse on both axes”. Registering the two directions a prediction could fail in is not enough when the mechanism can also simply fail to matter.

A second shape of the same failure, from a registration that offered two accounts and asked which one the data would pick. One said a correctly aimed quarter-sized mask would keep most of the learning of a half-sized one; the other said both quarters would cost more than the half and that aiming would buy nothing. The

measurement gave *both* effects, separately resolved and of comparable size: aim was worth $+0.0956 \pm 0.0063$ at matched size, and shrinking the mask at matched aim cost -0.0719 ± 0.0132 . **A dichotomy cannot be confirmed by a conjunction** — it can only be half-refused on each side, which is what the write-up had to say. The discipline that would have survived is to register the quantities and the direction each should move, not a menu of accounts to choose between.

2.12 Grep the results directory before building the experiment

Three times in one week the answer to a question this project was about to spend hours measuring was already in its own output. A symmetric account of leverage was refuted by a run on disk, not by a new one. A law relating cost to distance was refused by measuring a third point rather than by arguing about the first two. And a plan to re-measure a cost curve at the opposite end of the motor group — written up, committed, costed at about five hours — turned out to be an experiment that had run three weeks earlier, whose summary was filed in the note immediately adjacent to the one being corrected, and which the same day's work had cited for a different number.

The last of the three is the instructive one, because the obstacle was not that the result was hard to find. It was that the question had been *framed as a lead*, and a lead reads as something not yet done. The habit that catches it is mechanical rather than attentive: before building an arm, search the existing logs for the condition it would create — the position, the count, the target — and not for the name of the experiment that would create it. Arms are cheap to name and expensive to run, and the name is the one part that does not persist.

The same habit has a second form, which turned up within the hour of writing the first. Searching the archive *by arm name* finds every run of that arm and silently pools runs made under different conditions — §2.7 records what that cost. Both failures are one mistake: **treating a name as an identifier for a condition**. A name identifies an intention. The condition lives in the log beside it, and has to be read from there.

3. Credit assignment and interference

Teaching a second, conflicting vowel target destroys the first. The surviving fraction is 0.22. This section is what the project established about why, and about what does not fix it.

3.1 The wipe has a mechanism, and it is not a representational collision

Two lessons interfere while driving disjoint neuron groups, so the groups are not what they share. What they share is the reward: a single broadcast scalar nudges the bias of every neuron in the module, so teaching B random-walks A's setting. Measured directly, B restores the very neurons A suppressed. Orthogonal targets are held simultaneously at 0.84 against the conflicting pair's 0.22, which locates the failure in the conflict rather than in capacity.

3.2 It is an axis collision, not a distance effect

A second lesson that is silent about the first lesson's axis is free; one that speaks to the same axis is destructive, confirmed at +11.5 SE. Distance between the two targets is not the variable.

3.3 A perfect per-neuron credit mask does not remove the wipe

An oracle mask — the host tells the creature which lesson is live, and the reward is confined to the half of the motor group that lesson uses — retains *less* than the broadcast rule it was meant to repair: $+0.109 \pm 0.057$ against $+0.320 \pm 0.052$, a deficit of -0.211 ± 0.074 , negative at every gain floor. It also learns less, gaining $+0.196 \pm 0.015$ against $+0.362 \pm 0.033$. This retracts an earlier claim of the opposite sign, which had rested on the estimator described in §2.8.

3.4 The mask's cost decomposes additively, and content and position cost the same

On end-of-teaching error, paired over 32 creatures:

Step	Contrast	Effect
masking at all	no mask → mask aimed correctly	+0.211
content	correct aim → random aim	+0.103
position	random half → central half	+0.102
total	central mask vs no mask	$+0.417 \pm 0.044$ (9.5 SE)

The two increments are equal to within 0.001, and each is about half the base cost of masking at all. The direct position contrast — a central mask against a randomly aimed one, at matched content and matched mask fraction — is $+0.102 \pm 0.009$, 11.5 SE.

The ordering on learning is: no mask, then correct aim, then random aim, then central. That a mask excluding *both* extremes of the group costs as much again as mis-aiming one is consistent with the leverage geometry of §4.

3.5 The aim of the mask during the interfering lesson recovers half of what masking costs

Two arms that mask the first lesson identically and differ only in which half they confine the reward to during the *second* lesson. Both therefore learn the first lesson equally well — verified by end-of-teaching error identical to $+0.0000 \pm 0.0000$ over 32 creatures — so the comparison is clean.

Gain floor	Aim correct	Aim wrong	Difference	n
0.00	+0.1092	-0.0873	-0.1965 ± 0.0510 (3.9 SE)	32
0.05	+0.1291	-0.0815	-0.2107 ± 0.0509 (4.1 SE)	29
0.10	+0.1327	-0.0782	-0.2109 ± 0.0508 (4.2 SE)	27
0.20	+0.0598	-0.0934	-0.1532 ± 0.0478 (3.2 SE)	18

Retention as a paired, jackknifed ratio of sums, swept over a floor on per-creature learning. The effect holds at every floor that retains enough creatures, which is the criterion the experiment's own verdict gate applies.

Against a no-mask baseline of +0.320, a wrongly aimed mask sits -0.407 below it and correct aim recovers +0.197 — about 48%. The residual -0.211 is the part no amount of correct aim removes. So credit assignment here is not inert: what defeats the oracle is that the price of confining the reward exceeds what aiming it well buys back.

3.6 Position and content are not separable in a binary half-split

A structural result, and a negative one about the design rather than the creature. The mask is a contiguous half of the motor group and each lesson prefers exactly one half, so within a phase “the mask is in the wrong place” and “the mask selects the wrong neurons” are the same statement. No contrast built from end-aligned masks can separate them.

The only mask aligned with neither lesson — the central half, which leaves each lesson half of its preferred neurons on every trial — does separate them, and cannot answer the retention question for a second structural reason: it cuts learning by a factor of 29 (gain $+0.0125 \pm 0.0023$ against $+0.362$), so no creature clears even the lowest meaningful gain floor and retention has no denominator. **Varying position far enough to matter destroys the learning whose retention one would be measuring.** More creatures cannot fix this; it is not a power problem.

3.7 No mask in this family is cheap enough, and that closes the line by arithmetic

The two numbers above are a sum: confining the reward costs -0.211 ± 0.074 of retention and aiming it well recovers $+0.197$, so credit assignment here is outbid rather than inert. A cheaper mask was therefore the only thing that could change the sign, and §4.2 said where to look — most of the oracle’s half is low-potency ballast, so a small mask flush against the lever should keep the learning and pay less for it. Measured across every mask extent at the best available aim, on one genome, 32 creatures each:

Reward reaches	Aimed at	Learning gained
14 of 14 (no mask)	—	$+0.3618 \pm 0.0334$
7 of 14	the lever end	$+0.1964 \pm 0.0149$
4 of 14	the lever end	$+0.1245 \pm 0.0046$
7 of 14	a coin flip each trial	$+0.1022 \pm 0.0076$
4 of 14	the opposite end	$+0.0289 \pm 0.0054$
7 of 14, centred	neither end	$+0.0125 \pm 0.0023$

Aim at matched extent is $+0.0956 \pm 0.0063$ (15.2 SE) paired on seed; extent at matched aim is -0.0719 ± 0.0132 (5.4 SE).

Learning rises monotonically with how much of the group the reward can reach, and the cheapest mask is no mask. The best-aimed quarter keeps 63% of the half’s learning on 57% of its neurons, so potency does not make a mask cheap — it makes a mask less dear than the same mask pointed the other way. Confinement is what costs; aim reduces that cost and never reverses it. So there is no mask in this family that credit assignment can afford, and the line closes on arithmetic rather than on anything about the creature.

4. Leverage and the geometry of the motor pool

4.1 The lesson silences the wrong side

Measured rather than fitted: a low formant target is reached by *suppressing* the upper half of the group, and a high target by suppressing the lower half — mirror-image profiles, both suppression. An earlier account in which the lesson raised the half it needed is a decisive reversal, and the reversal is what the rest of this section is an account of.

4.2 The readout is a centroid, so output potency is *signed* distance from it

The formant the host reads off a motor group is not a mean of rates. It is a rate-weighted centroid on a labelled line: neuron i carries a fixed position $p_i = (i - \text{begin} + \frac{1}{2}) / n$ and the output is $\sum(r_i p_i) / \sum r_i$. Differentiating gives the influence of one neuron on behaviour with no free parameter:

$$\partial F / \partial r_i = (p_i - c) / \sum r_j \quad c = \text{the output currently produced}$$

So potency is distance from the centroid the group is *currently* at — zero for a neuron sitting on it, largest at the ends, and state-dependent, moving as c moves. That immediately explains a documented out-of-sample failure: a leverage law fitted to three single-neuron costs with one free parameter reproduced them to 1.68 percentage points and then failed two half-block costs. The fit was a function of *position alone*, which is adequate while c barely moves — as it does not between single-neuron blocks — and wrong once a whole half of the group is blocked and c moves a long way. Influence is not a property of a neuron; it is a property of a neuron and the current state.

The magnitude $|p_i - c|$ is nevertheless refuted, by a run that was already on disk when the derivative was written down. One neuron blocked at three positions, 36 creatures: the top ($p = 0.964$) costs 30% of the lesson at -5.8 SE, the middle (0.464) costs 7% at -1.1 SE, the bottom (0.036) costs 9% at -1.4 SE. Middle and bottom are both null and indistinguishable from each other, and no single centroid makes a symmetric magnitude produce that: at $c = 0.093$ it requires the middle to cost about six times the bottom, and at $c \rightarrow 0.5$ it requires the bottom to cost like the top.

Potency is $(p_i - c)$ taken with the sign of the required direction of travel. A lesson that must lower the centroid can only do so by suppressing high-position neurons, so the far end *on that side* carries the lever; the centroid itself has no lever in either direction; and the opposite end is on the wrong side, because suppressing it would move the output the wrong way, so the lesson was never going to use it.

The qualitative half of this was written before the run that measured it, in the pre-registration of that experiment: index is position, a weighted mean is least sensitive near its own centre, and the end the lesson travels toward should matter more. What is new is the derivative itself, the state-dependence, and the sign — and the first arithmetic written from it, the symmetric one, was wrong.

4.3 Exactly one neuron of fourteen is expensive, and which one depends on the target

The prediction the signed account makes and a symmetric one cannot: changing where the lesson is going must *swap* which end is dear, not grade it. Three lesson targets, 36 creatures each, each arm paired on seed against its own unblocked control:

Target	p	Travel	Lever	Cost of the lever	Middle	Far end
320 Hz	0.093	0.407	top	30.0% ± 5.1%	7%	9%
475 Hz	0.300	0.200	top	25.0% ± 5.9%	6%	5%
850 Hz	0.800	0.300	bottom	22.1% ± 4.4%	4%	3%

Cost as a percentage of the lesson learned by the unblocked control at the same target. Both end columns are the single neuron at that end of the group of fourteen.

Three targets; three times the lever is the single neuron at the end the lesson must *suppress*; three times the middle is free. Every dear arm clears 4 SE and every cheap arm sits under 1.5 SE at comparable standard errors, so the cheapness is measured rather than merely unresolved. The two low targets and the high one place the lever at opposite ends of the same group, which is the swap rather than a gradient.

Both later runs carried the original target's four arms unchanged, and its headline cost reproduced to four decimals (-0.0503 ± 0.0086 , -5.8 SE) in each. The runs are therefore self-validating rather than resting on a cross-run comparison.

4.4 The size of the lever is not a function of how far the lesson travels

A second quantity looked as though it fell out of the same account and does not. Across the first two targets, the distance the centroid had to travel and the cost of the lever stood in ratios of 1.356 and 1.357 — agreement to 0.1% with nothing fitted. The third target refused it outright, and the costs are non-monotone in travel, so a cost proportional to distance cannot produce the ordering at all. §2.10 records the inference lesson. What survives is the direction result of §4.3 and an explicit refusal to fit the magnitude: the end the lesson must suppress carries the lever, the middle never does, and what sets the size of that lever is open.

4.5 Position decides whether extent costs anything at all

The same two axes, measured independently in the other family and on one reference lesson of $+0.1676 \pm 0.0073$, 36 creatures. Here neurons are *frozen* rather than being the only ones rewarded, so the axis runs the other way: the group is fourteen, the lesson is a low target, and the ladder is built from each end.

Frozen	At the lever end	At the opposite end
1 of 14	-0.0503 ± 0.0086 30.0% —	—
2 of 14	-0.0599 ± 0.0088 35.7%	-0.0144 ± 0.0099 8.6%
4 of 14	-0.0957 ± 0.0091 57.1%	-0.0171 ± 0.0095 10.2%
7 of 14	-0.0889 ± 0.0080 53.1%	-0.0153 ± 0.0102 9.1%

Cost of freezing, paired on seed against the unfrozen control, as a fraction of that control's own lesson. The two ends at seven neurons each are complementary halves of the group.

Read down each column. **At the lever end extent costs: four neurons against one is -0.0455 ± 0.0075 , 6.1 SE. At the opposite end it costs nothing at any size —** four against two is -0.0028 ± 0.0092 (0.3 SE), seven against four is $+0.0019 \pm 0.0087$ (0.2 SE), seven against two is -0.0009 ± 0.0086 (0.1 SE). Three steps

under half a standard error, at standard errors small enough to have resolved a step the size of the one the other ladder shows, so the flatness is measured rather than unresolved. At matched extent the two ends differ by 8.6 to 8.7 SE.

So position is not one cost among two: **it decides whether the other axis exists**. Freezing more of the end the lesson must suppress costs more every time until it saturates; freezing more of the end the lesson was never going to use costs nothing, however much of it goes. §3.7 reaches the same two-axis statement from the reward-mask family by a different manipulation, and the two were measured three weeks apart.

The lever end also *saturates*, and that is the second column’s own result: seven neurons against four is $+0.0068 \pm 0.0068$, 1.0 SE. Once the top four of fourteen are frozen there is nothing further to lose, which is why the costs at four and at seven are the same number.

This corrects a claim of this project’s own, which had been recorded as “the block cost is flat in the fraction blocked, so there is no mask size that wins”. That rests on exactly the two points above the saturation knee — four and seven at the lever end — while the one-to-four segment of the same ladder rises from 30% to 57% at 6.1 SE. The flatness is real and local: it is the saturation of one ladder above four of fourteen, together with the whole of the other ladder sitting at zero. It is not a property of the fraction blocked, and it was drawn from a subset of a curve the project had already measured in full.

4.6 What a restriction costs has two axes: the potency of what it restricts, and how much

Stated across every plasticity restriction the project has built, with the columns measured in separate experiments:

Restriction	Selects by	Effect on learning	Effect on forgetting
Astrocytic activity gate	activity	−81%	−19%
Per-neuron reward mask, aimed correctly	position	−46%	−30%
Per-neuron reward mask, central	position	−97% (29×)	undefined
Consolidation brake on high-traffic edges	traffic	null	null

The brake is the informative row. It binds hard — mean per-edge learning rate 0.6708 against 1.0000 unbraked, reproducing an independent earlier measurement of 0.654 — and it changes neither retention (−0.6 SE, flat across every gain floor) nor learning. A mechanism that restricts a third of all plasticity and costs nothing is not a weak version of a mask that restricts half and costs 46%: it selects a *different set*. Traffic is uncorrelated with position in the group, so a traffic-selected restriction spreads across the centroid and leaves the ends free, while a position-selected one of the same size can remove the lever entirely. **Extent alone buys nothing: a restriction whose potency is zero is free however wide it is.**

Written here first as “size is the wrong axis, potency is the right one”, and that is half wrong. §3.7 measured extent at matched aim and it costs -0.0719 ± 0.0132 , 5.4 SE, and §4.5 has the same two axes in the freezing family at 6.1 SE against 0.1 SE. The statement that survives all three is: **potency decides whether extent costs anything, and given nonzero potency extent then costs until it saturates**. The brake is the potency-zero end of it and the mask ladder is the other; neither is the whole.

The first three rows are a quantitative form of the stability–plasticity dilemma in one system — every restriction cut learning more than it cut forgetting — and the fourth is not an exception to the ordering so much as a case where the dilemma never engaged. This project predicted the opposite for the brake before

running it, which §2.11 records.

4.7 The larynx reads nine scalars

The motor interface is nine scalar groups, and the object identity being taught sits at the null in all eighteen vocal measurements taken. Reordering the vocal population by its own response difference lifts discriminability, so the information is present in the slice and absent from the readout.

5. The voice and the body

5.1 The voice drones

There is no syllabic rhythm in the produced voice. A +9.99 dB modulation peak that indicated otherwise was a fit on the knee of a low-pass filter, not a spectral feature of the signal. Relatedly, a 3 Hz ring in the voice is the product of two filters — a fast rate constant and a gate smoothing term — and may be no oscillator at all; it scales threefold with self-hearing, and a deaf creature still rings at +3.1 SE, so self-audition is a gain on it rather than its cause.

5.2 The frame is a body part

A single finding that explains five separate refusals at once: the apparent syllabic frame is a resonant jaw at about 4.5 Hz, a property of the articulator rather than a central clock. Three independent routes from a frame to content are closed by it, and a jaw-gate mechanism that appeared to pass both halves of its test turned out to have a mis-set anchor, with the creature's own silences outperforming the imposed gate.

5.3 A blind controller beats feedback

A controller with no access to the sensed formant outperforms one with it: reach is immune to loop delay and a 905 ms loop is not. The precision advantage claimed for feedback was cleared by a sensorless controller, and an uninformative feature leaves only the intercept, which is the blind controller in disguise. Altered auditory feedback moves production by +0.01.

5.4 Intrinsic plasticity is the backlash

Formant actuator hysteresis measures 107 Hz with intrinsic plasticity enabled and 2.0 Hz with it off: the plant moves under its own controller. The backlash is per-neuron — pooling the regulator takes hysteresis from 107 Hz to 7 Hz and gives back 30 Hz of reach.

5.5 The ear hears its own formant and does not use it

The self-produced first formant is present in the auditory representation (+0.124 and +0.114 on two seed families) and unused by the motor system. The ear's estimate is well described by a fixed attenuation plus noise, with the attenuation free and the noise 56 Hz input-referred; the signal is localised, one frequency bin of 32 carrying 92% of the available benefit, which is destroyed by pooling. So the failure of closed-loop control here is a missing *mapping*, not a missing signal.

6. Representation and capacity

6.1 Vocabulary does not collapse, and nine words do not cost naming

Advantage over chance *rises* with vocabulary size rather than falling. At nine words against two six-word arrangements, total naming excess is conserved (-0.012 ± 0.019 and $+0.020 \pm 0.018$, opposite-sign nulls at $n = 36$). There is one real cost and it is not naming: delivered formant separation falls by 11.8 ± 3.5 Hz at matched trials per word, 3.3 SE. The shape of the word arrangement does not matter (-0.032 ± 0.027), so there is no preferred axis per creature.

6.2 The ceiling is an exponent

Delivered formant change scales with the aligned component of the learned bias as a power law with exponent about 0.61. Moving delivery by a factor requires moving the bias by a larger factor, which closes the naming line quantitatively rather than by argument: a commitment brake moved the aligned bias $1.07\times$ where $6.2\times$ was needed.

6.3 The derived index gap has no surviving cause

The creature derives its own context from an auditory rate average, and the resulting index is reliably *worse* when the teaching target tracks the heard word than when it is fixed, on 18 of 18 creatures — a gap of -0.056 ± 0.009 , 6.0 SE. Three candidate causes were tested and all three are dead: gate corruption has the wrong sign; silencing self-audition does not close the gap but *widens* it by 0.037 ± 0.015 , under a test with power to see a closure; and a drift account is excluded both by its own point prediction, 3.9 SE away, and by the learning arm decaying least. The gap is real and unexplained.

Two side results survived that line. Deafening improves the index in the control arm only ($+0.030 \pm 0.004$, 7.1 SE), and makes the tracking arm's partition more lopsided ($+0.066 \pm 0.018$, 3.7 SE). Separately, the derived context carrier is refuted directly: the index never separates the two lessons (0.23 against 0.23), while the auditory feature it is built from carries the word at 97.8% on held-out trials where the reward lands. The feature is fine; the extraction is not.

6.4 Two percent of a percent of the ticks does the work

0.2% of ticks account for the learning, and no schedule reproduces it at 3.3 SE over 32 creatures. The result carries no lesson about *which* ticks, which is stated here because it is the kind of finding that invites one.

7. Negative results that closed lines

Listed because the project's cost is dominated by them and because each removed a plausible mechanism from consideration.

Mechanism	Outcome
Oscillations as a rate-to-phase code	Built; spike-timing plasticity could not read it
Hippocampal pattern separation	Does not separate

Exuberant synaptogenesis	Costs 0.229; competition a further 0.051. Weight is not information
Dendritic error microcircuit	Settles and cancels nothing
Plateau-gated plasticity	Attenuates instead of selecting
Two-stage bandpass	Never appears on three seeds
Topographic tract	Oracle predicted 471 Hz; measured 1.00×
Region-shaped reward targets	A satisfied lesson still does not protect the first
Two-pool differential readout	Premise fails: only $\rho = 0.101$ of variance is shared
Adaptation current as a clock	Peak frequency follows its time constant at exponent 0.10 where 1.0 is required
Prototype freezing	Refused; it is the stability–plasticity dilemma
Replay / interleaving as shipped	Switching it off costs 0.3 SE
Reward that sees target distance	Worse, by 9.4 SE
Protocol axes (duration, blocking)	28× longer is no help; blocked teaching is harmful
Astrocytic activity gate on plasticity	Cuts learning 81% and forgetting 19%, on a monotonic dose–response
Consolidation brake on high-traffic edges	Binds hard (0.671 of the unbraked rate) and changes neither retention nor learning

8. Provenance

Every number above comes from a run of the project’s own experiment binary, gated before execution on --experiment verify reporting PASS with the pinned hash ad96f882becbee92 and a passing pinned-hash check. Runs that failed the gate are discarded rather than reported.

Sample sizes are creatures, not trials: a creature is one genome with one seed, grown and taught from birth. Where a paired statistic is quoted, creatures are matched by seed across arms, which is licensed by the per-job seed being a function of the repetition index alone and independent of the arm — verified by a shared baseline arm reproducing bit-identically across three separate multi-hour runs, and by a direct cross-run check of that arm against itself returning $+0.0000 \pm 0.0000$ over 32 creatures.

Retention figures are a ratio of sums with a leave-one-out jackknife standard error, swept over a floor on per-creature learning, and are reported only where the effect holds at every floor retaining at least three creatures in both arms. Section 2.8 explains why the more obvious estimator is not used.

Run logs, genome files and the experiment source are in the project repository. The findings here are recorded at greater length, with the reasoning that produced them and the criteria that would refuse them, in the project’s README.md and its memory store.

Appendix A. Excluded by the criterion

These project results are omitted from the body because a published result says or frames them first. They are listed so that the inclusion criterion of §1 can be checked.

One entry needs its line drawn explicitly. The idea that behaviour depends on some directions in a population code and not others is Kaufman and colleagues', and it is excluded. What §4.2 claims is narrower and is included: in this system the map from population activity to behaviour is *analytically known rather than fitted*, so the potent direction can be differentiated instead of estimated, it is state-dependent, and its sign — not its magnitude — is what the measurements require.

Project result

External antecedent

Interference as a failure of contextual inference; a per-neuron mask should therefore help	Heald, Lengyel & Wolpert, COIN
Gating helps and is insufficient alone; stabilisation is the missing half	Masse, Grant & Freedman (2018)
Importance-weighted protection of earlier weights as the remedy	Kirkpatrick et al. (2017); Zenke, Poole & Ganguli (2017)
Node perturbation as the motor learning rule	Fiete, Fee & Seung
Sequence generation needs competition, not spike timing alone	Fiete, Senn, Wang & Hahnloser (2010)
Prototype freezing as a dilemma rather than a bug	Grossberg
A uniform cluster prior implicit in a competitive learning rule	MacQueen; Assran et al. (2023)
Ratio-of-sums estimation and the jackknife	Cochran; Quenouille; Tukey
Per-context reward prediction error having nothing to remove	Gadagkar et al. (2016)
Behaviour depending on output-potent and not output-null dimensions of a population code	Kaufman, Churchland, Ryu & Shenoy (2014)
Activity-dependent myelination as a consolidation mechanism, and its post-training window	The oligodendroglia literature, cited by title and DOI
An astrocyte-gated multi-timescale rule as a remedy for catastrophic forgetting	Astrocyte-gated multi-timescale plasticity (2026)
Contexts lying on a continuum rather than being discrete slots	Heald, Lengyel & Wolpert (2023)